

In [2]:

```
import numpy as np
import pandas as pd
import seaborn as sb
import matplotlib.pyplot as pp
```

In [3]:

```
df1 = pd.read_csv(r"C:\Users\user\Desktop\c10\madrid_2018.csv")
df = df1.head(1000)
df
```

Out[3]:

	date	BEN	CH4	CO	EBE	NMHC	NO	NO_2	NOx	O_3	PM10	PM25	SO_2	1
0	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	29.0	31.0	NaN	NaN	NaN	2.0	1
1	2018-03-01 01:00:00	0.5	1.39	0.3	0.2	0.02	6.0	40.0	49.0	52.0	5.0	4.0	3.0	1
2	2018-03-01 01:00:00	0.4	NaN	NaN	0.2	NaN	4.0	41.0	47.0	NaN	NaN	NaN	NaN	1
3	2018-03-01 01:00:00	NaN	NaN	0.3	NaN	NaN	1.0	35.0	37.0	54.0	NaN	NaN	NaN	1
4	2018-03-01 01:00:00	NaN	NaN	NaN	NaN	NaN	1.0	27.0	29.0	49.0	NaN	NaN	3.0	1
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
995	2018-03-02 18:00:00	NaN	NaN	0.4	NaN	NaN	13.0	62.0	83.0	31.0	NaN	NaN	NaN	1
996	2018-03-02 18:00:00	NaN	NaN	NaN	NaN	NaN	19.0	70.0	99.0	NaN	9.0	NaN	4.0	1
997	2018-03-02 18:00:00	NaN	NaN	NaN	NaN	NaN	42.0	88.0	152.0	NaN	3.0	2.0	NaN	1
998	2018-03-02 18:00:00	NaN	NaN	NaN	NaN	NaN	20.0	69.0	100.0	NaN	11.0	10.0	NaN	1
999	2018-03-02 18:00:00	NaN	NaN	NaN	NaN	NaN	19.0	76.0	105.0	14.0	NaN	NaN	NaN	1

1000 rows × 16 columns



In [4]:

```
df=df.dropna()
```

In [5]:

```
df.columns
```

Out[5]:

```
Index(['date', 'BEN', 'CH4', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'NOx', 'O_3',
      'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 83 entries, 1 to 990
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        83 non-null    object
1   BEN         83 non-null    float64
2   CH4         83 non-null    float64
3   CO          83 non-null    float64
4   EBE         83 non-null    float64
5   NMHC        83 non-null    float64
6   NO          83 non-null    float64
7   NO_2        83 non-null    float64
8   NOx         83 non-null    float64
9   O_3         83 non-null    float64
10  PM10        83 non-null    float64
11  PM25        83 non-null    float64
12  SO_2        83 non-null    float64
13  TCH         83 non-null    float64
14  TOL         83 non-null    float64
15  station     83 non-null    int64
dtypes: float64(14), int64(1), object(1)
memory usage: 11.0+ KB
```

In [7]:

```
data=df[['CO' , 'station']]
data
```

Out[7]:

	CO	station
1	0.3	28079008
6	0.2	28079024
25	0.2	28079008
30	0.2	28079024
49	0.2	28079008
...	...	...
942	0.3	28079024
961	0.6	28079008
966	0.2	28079024
985	0.5	28079008
990	0.2	28079024


83 rows × 2 columns

In [8]:

```
data.plot.line(subplots=True)
```

Out[8]:

array([&lt;AxesSubplot:~&gt;, &lt;AxesSubplot:~&gt;], dtype=object)




In [9]:

```
data.plot.line()
```

Out[9]:

&lt;AxesSubplot:&gt;



In [10]:


```
x = data[0:100]
```

In [11]:

```
x.plot.bar()
```

Out[11]:

&lt;AxesSubplot:&gt;




In [12]:

```
data.plot.hist()
```

Out[12]:

&lt;AxesSubplot:ylabel='Frequency'&gt;




In [13]:

```
data.plot.area()
```

Out[13]:

&lt;AxesSubplot:&gt;




In [14]:

```
data.plot.box()
```

Out[14]:

<AxesSubplot:>



In [15]:

```
x.plot.pie(y='station' )
```

Out[15]:


```
<AxesSubplot:ylabel='station'>
```

In [16]:

```
data.plot.scatter(x='CO' ,y='station')
```

Out[16]:

<AxesSubplot:xlabel='CO', ylabel='station'>





In [17]:

df.info()

&lt;class 'pandas.core.frame.DataFrame'&gt;

Int64Index: 83 entries, 1 to 990

Data columns (total 16 columns):

# Column Non-Null Count Dtype

```

---
0 date      83 non-null object
1 BEN      83 non-null float64
2 CH4      83 non-null float64
3 CO       83 non-null float64
4 EBE      83 non-null float64
5 NMHC     83 non-null float64
6 NO       83 non-null float64
7 NO_2     83 non-null float64
8 NOx      83 non-null float64
9 O_3      83 non-null float64
10 PM10    83 non-null float64
11 PM25    83 non-null float64
12 SO_2    83 non-null float64
13 TCH     83 non-null float64
14 TOL     83 non-null float64

```

In [18]:

df.describe()

Out[18]:

	BEN	CH4	CO	EBE	NMHC	NO	NO_2	TCH
count	83.000000	83.000000	83.000000	83.000000	83.000000	83.000000	83.000000	83.000000
mean	0.589157	1.275301	0.265060	0.185542	0.044096	12.783133	36.843373	56.385
std	0.276305	0.169010	0.109804	0.123103	0.020184	15.798308	27.034974	50.024
min	0.200000	1.080000	0.100000	0.100000	0.010000	1.000000	2.000000	3.000
25%	0.400000	1.100000	0.200000	0.100000	0.030000	1.000000	14.000000	17.000
50%	0.600000	1.350000	0.299000	0.100000	0.040000	4.000000	34.000000	40.000
75%	0.700000	1.400000	0.300000	0.200000	0.050000	19.500000	55.000000	89.000
max	1.600000	1.720000	0.600000	0.600000	0.110000	71.000000	102.000000	208.000

In [20]:

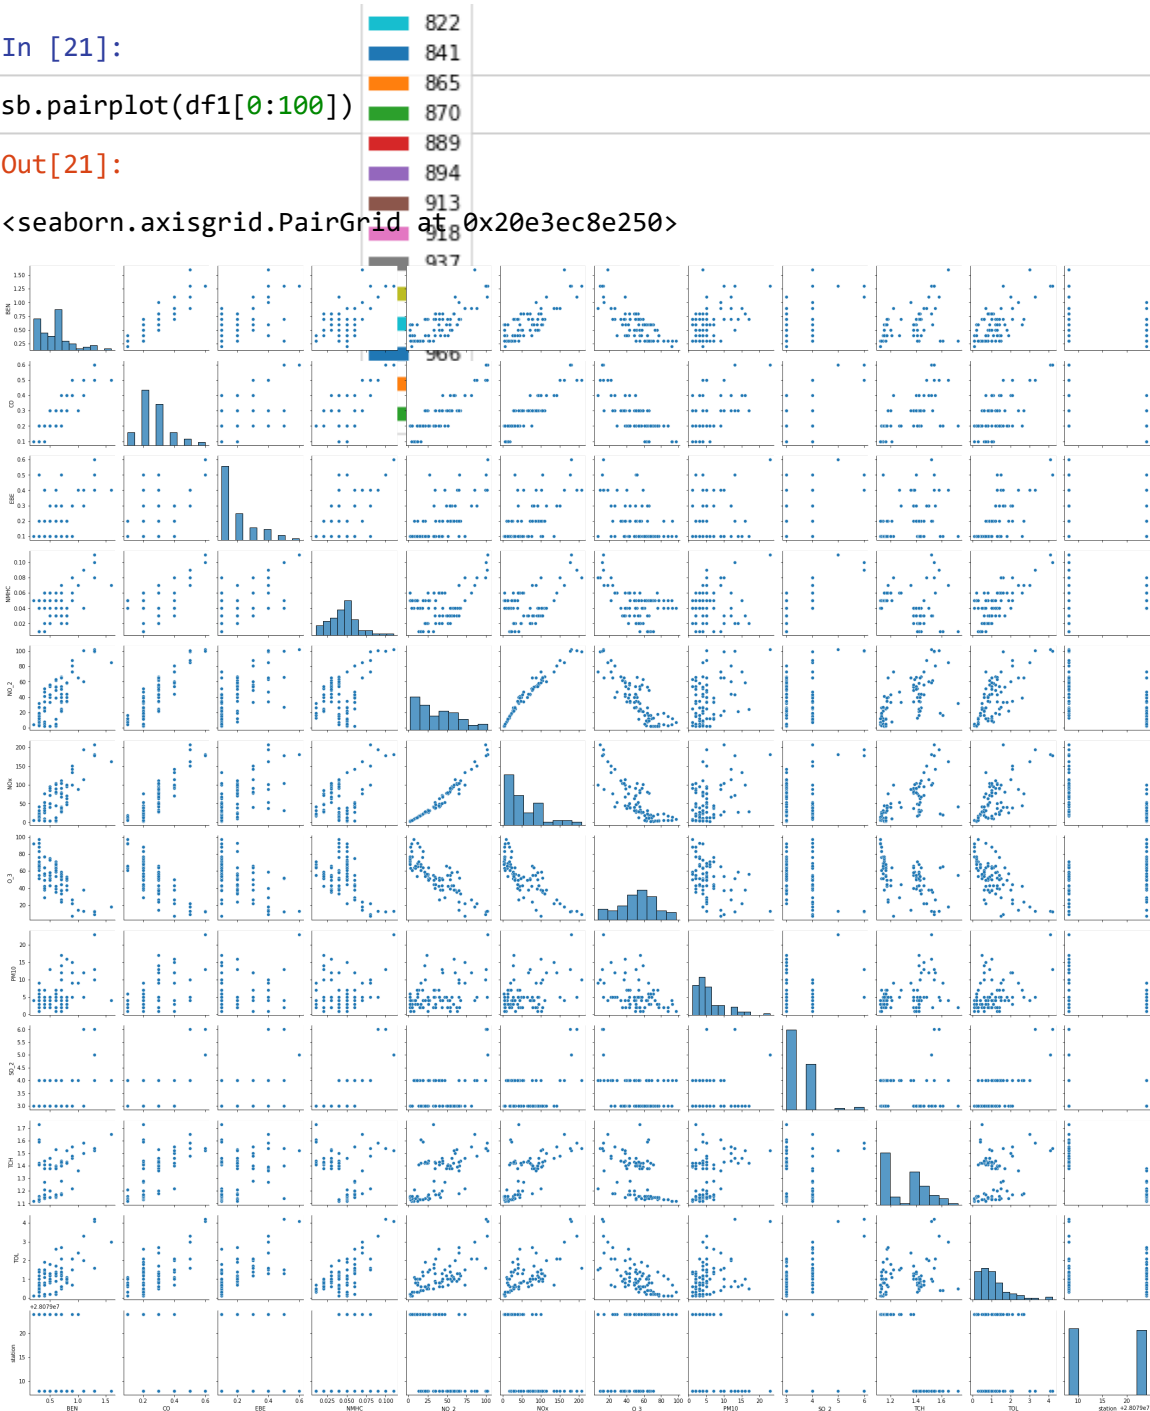
```
df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NOx', 'O_3',
        'PM10', 'SO_2', 'TCH', 'TOL', 'station']]
```

In [21]:

```
sb.pairplot(df1[0:100])
```

Out[21]:

<seaborn.axisgrid.PairGrid at 0x20e3ec8e250>



In [22]:


```
sb.distplot(df1['station'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:255  
7: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

Out[22]:

<AxesSubplot:xlabel='station', ylabel='Density'>



In [ ]:

```
sb.heatmap(df1.corr())
```

In [23]:

```
x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NOx', 'O_3',  
      'PM10', 'SO_2', 'TCH', 'TOL']]  
y=df['station']
```

In [24]:

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

In [25]:

```
from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

Out[25]:

LinearRegression()

In [26]:

```
lr.intercept_
```

Out[26]:

28079031.416664694

In [27]:

```
coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff
```

Out[27]:


	Co-efficient
<b>BEN</b>	2.850632
<b>CO</b>	-33.764002
<b>EBE</b>	2.894612
<b>NMHC</b>	263.064382
<b>NO_2</b>	0.093230
<b>NOx</b>	-0.121530
<b>O_3</b>	-0.043294
<b>PM10</b>	0.019408
<b>SO_2</b>	-0.305820
<b>TCH</b>	-10.166375
<b>TOL</b>	-0.616478

In [28]:

```
prediction =lr.predict(x_test)
pp.scatter(y_test,prediction)
```

Out[28]:

&lt;matplotlib.collections.PathCollection at 0x20e47e3f070&gt;



In [29]:

```
lr.score(x_test,y_test)
```

Out[29]:

0.8593006551325477

In [30]:

```
lr.score(x_train,y_train)
```

Out[30]:

0.9685777662792087

In [31]:

```
from sklearn.linear_model import Ridge,Lasso
```

In [32]:

```
r=Ridge(alpha=10)  
r.fit(x_train,y_train)
```

Out[32]:

Ridge(alpha=10)

In [33]:

```
r.score(x_test,y_test)
```

Out[33]:

0.6843527368959668

In [34]:

```
r.score(x_train,y_train)
```

Out[34]:

0.6805566756117155

In [35]:

```
l=Lasso(alpha=10)  
l.fit(x_train,y_train)
```

Out[35]:

Lasso(alpha=10)

In [36]:

```
l.score(x_train,y_train)
```

Out[36]:

0.43550728759496293

In [37]:

```
l.score(x_test,y_test)
```

Out[37]:

0.3738345569498982

In [38]:

```
from sklearn.linear_model import ElasticNet
e=ElasticNet()
e.fit(x_train,y_train)
```

Out[38]:

ElasticNet()

In [39]:

```
e.coef_
```

Out[39]:

```
array([ 0.          , -0.          ,  0.          ,  0.          , -0.05044841,
        -0.14127213, -0.16530384, -0.21242968,  1.74123806, -0.0257283 ,
         0.640385   ])
```

In [40]:

```
e.intercept_
```

Out[40]:

28079028.52312789

In [41]:

```
prediction=e.predict(x_test)
```

In [42]:

```
e.score(x_test,y_test)
```

Out[42]:

0.5566407016409418

In [43]:

```
from sklearn import metrics
```

In [44]:

```
print(metrics.mean_squared_error(y_test,prediction))
```

27.239995291180534

In [45]:

```
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

5.219194889174051

In [46]:

```
print(metrics.mean_absolute_error(y_test,prediction))
```

4.456924369931221

In [47]:

```
from sklearn.linear_model import LogisticRegression
```

In [49]:

```
feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'NOx', 'O_3',  
                  'PM10', 'SO_2', 'TCH', 'TOL']]  
target_vector=df[ 'station']
```

In [50]:

```
feature_matrix.shape
```

Out[50]:

(83, 11)

In [51]:

```
target_vector.shape
```

Out[51]:

(83,)

In [52]:

```
from sklearn.preprocessing import StandardScaler
```

In [53]:

```
fs=StandardScaler().fit_transform(feature_matrix)
```

In [54]:

```
logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

Out[54]:

LogisticRegression(max\_iter=10000)

In [60]:

```
observation=[[1,2,3,4,5,6,7,8,9,10,11]]
```

In [61]:

```
prediction=logr.predict(observation)  
print(prediction)
```

[28079008]

In [62]:

```
logr.classes_
```

Out[62]:

```
array([28079008, 28079024], dtype=int64)
```

In [63]:

```
logr.score(fs,target_vector)
```

Out[63]:

```
1.0
```

In [64]:

```
logr.predict_proba(observation)[0][0]
```

Out[64]:

```
0.9999999988636115
```

In [65]:

```
logr.predict_proba(observation)
```

Out[65]:

```
array([[9.99999999e-01, 1.13638856e-09]])
```

In [66]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [67]:

```
rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

Out[67]:

```
RandomForestClassifier()
```

In [68]:

```
parameters={'max_depth':[1,2,3,4,5],  
            'min_samples_leaf':[5,10,15,20,25],  
            'n_estimators':[10,20,30,40,50]  
}
```



In [69]:

```
from sklearn.model_selection import GridSearchCV
grid_search = GridSearchCV(estimator=rfc, param_grid=parameters, cv=2, scoring="accuracy")
grid_search.fit(x_train, y_train)
```

Out[69]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),
             param_grid={'max_depth': [1, 2, 3, 4, 5],
                         'min_samples_leaf': [5, 10, 15, 20, 25],
                         'n_estimators': [10, 20, 30, 40, 50]},
             scoring='accuracy')
```

In [70]:

```
grid_search.best_score_
```

Out[70]:

```
0.9827586206896552
```

In [71]:

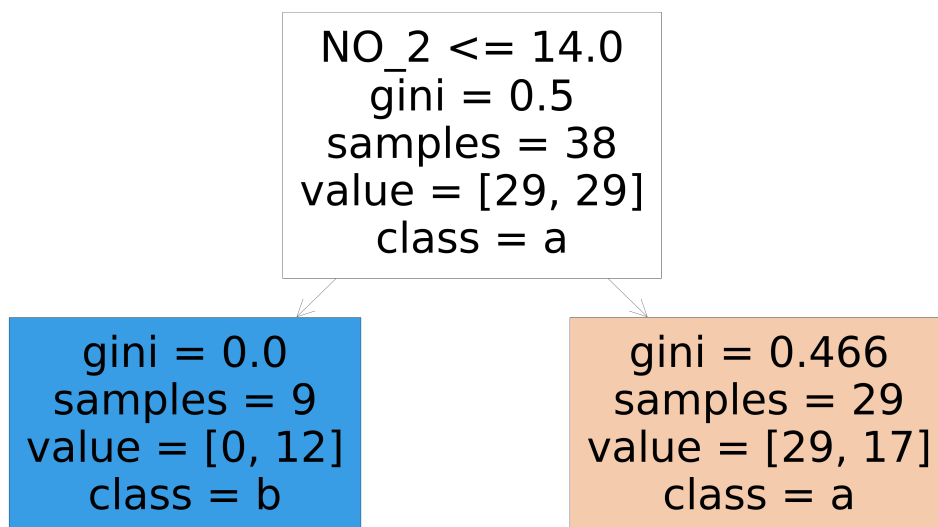
```
rfc_best=grid_search.best_estimator_
```

In [72]:

```
from sklearn.tree import plot_tree
pp.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['a', 'b', 'c', 'd'], f
```

Out[72]:

```
[Text(2232.0, 1630.8000000000002, 'NO_2 <= 14.0\n gini = 0.5\n samples = 38\n \nvalue = [29, 29]\n \nclass = a'),
 Text(1116.0, 543.5999999999999, 'gini = 0.0\n samples = 9\n \nvalue = [0, 12]\n \nclass = b'),
 Text(3348.0, 543.5999999999999, 'gini = 0.466\n samples = 29\n \nvalue = [29, 17]\n \nclass = a')]
```



# random forest is best suitable for this dataset

In [ ]: